

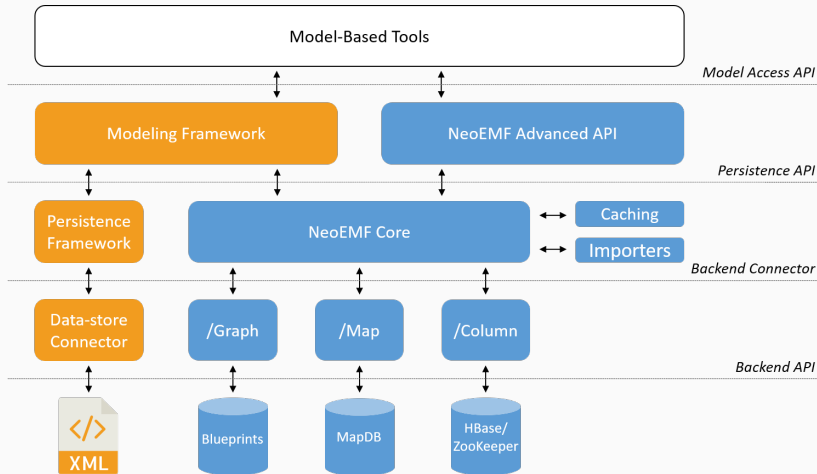
Taming Large Models with Hawk and NeoEMF

A. García-Domínguez, D. S. Kolovos, K. Barmpis, G. Daniel, G. Sunyé
MoDELS'2018, 14–19 October 2018

NeoEMF

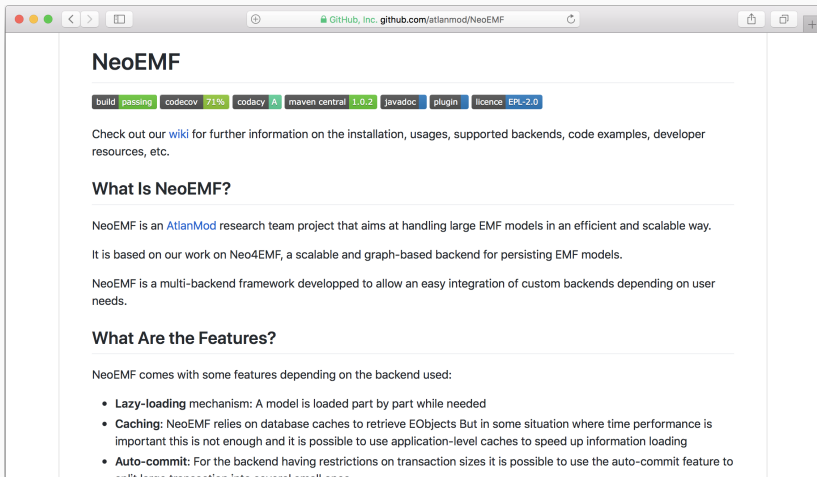
- Handle large models with task-specific databases
- Lazy-loading
- Compliant with the EMF API
 - Easy to integrate in existing applications
 - EMF-Compatible code generation
- Advanced caching (& prefetching strategies)
- Efficient XMI importer

NeoEMF: Architecture



- **NeoEMF/Graph**
 - Efficient model traversal using rich query language
 - Mogwai framework (OCL to Gremlin translation)
- **NeoEMF/Map**
 - Fast access to atomic operations
 - Designed for EMF-API calls
- **NeoEMF/Column**
 - Transparent model distribution
 - Concurrent read/write
 - Distributed model transformations (ATL-MR)

NeoEMF: project website



- <https://github.com/atlanmod/NeoEMF>
- Open source project under the Eclipse Public License 2.0

NeoEMF: initialise a new resource

1. Register the Persistence Backend Factory.
2. Create a ResourceSet and register the PersistentResourceFactory
3. Create a new URI to locate a file-based resource.
4. Create the resource.

```
PersistenceBackendFactoryRegistry.register(  
    BlueprintsURI.SCHEME,  
    BlueprintsPersistenceBackendFactory.getInstance());  
  
ResourceSet resourceSet = new ResourceSetImpl();  
resourceSet.getResourceFactoryRegistry().getProtocolToFactoryMap()  
    .put(BlueprintsURI.SCHEME,  
        PersistentResourceFactory.getInstance());  
  
URI uri = BlueprintsURI.createFileURI(new File("<db_path>"));  
  
Resource resource = resourceSet.createResource(uri);  
  
// EMF resource stored in an in-memory Blueprints graph
```

NeoEMF: persist a resource

1. Create a new option builder (backend-specific).
2. Save the resource.
3. Manipulate the resource by accessing the local database

```
Map<String, Object> options = BlueprintsNeo4jOptionsBuilder.newBuilder()  
    .weakCache().autocommit().cacheSizes()
```

```
resource.save(options);  
// Resource saved in a local Neo4j database  
resource.getContents()  
// [...]
```

NeoEMF: modify an existing resource

1. Load resource using the same option builder.
 2. Navigate its content and perform update operations
 3. Save the resource (automatically done with *autocommit* option)
-

```
Map<String, Object> options = BlueprintsNeo4jOptionsBuilder.newBuilder()  
    .weakCache().autocommit().cacheSizes()
```

```
URI uri = BlueprintsURI.createFileURI(new File("<db_path>"));
```

```
resourceSet.createResource(uri)  
resource.load(options);
```

```
// Model manipulation operation (complete EMF API support)  
MyClass myClass = (MyClass) resource.getContents().get(0);  
myClass.setName("NewName");
```

```
// Save the modification in the local Neo4j database  
resource.save(config.asMap());
```


Demo time!

Let's import a Java model, save it in Neo4j
and MapDB, and query the database.

Mogwai

Mogwai: motivation

- NeoEMF improves model scalability, but ...



Under the Hood

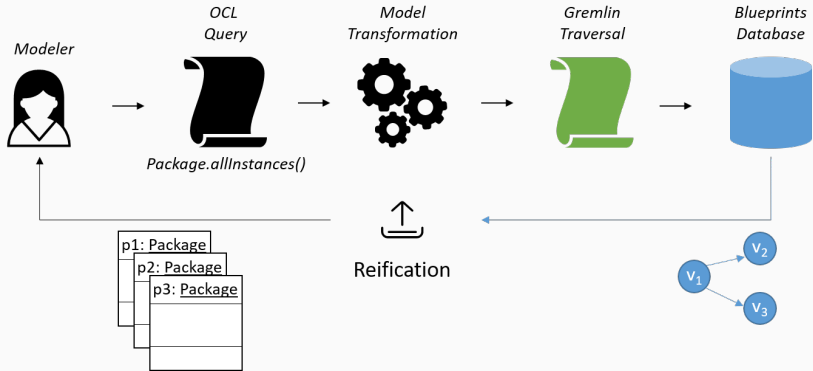
`Package.allInstances().name`

→
`get(p1)`
`get(p1,name)`
...
`get(pn)`
`get(pn,name)`

- Low-level model handling APIs
 - Not aligned with the database capabilities
- Fragmented queries on the database
 - Not efficient
 - Remote databases
- Intermediate object reification
 - Memory consumption
 - Execution time overhead

- Database queries are efficient but
 - Modern persistence frameworks typically rely on NoSQL databases
 - Multiple query languages
 - Multiple data representations
 - Low-level queries are hard to understand and maintain
 - Modeling expertise vs. Database expertise
- Solution: generate them!

Mogwai: architecture [DSC16]



- Generate graph database queries from OCL expressions
- Bypass modelling framework API
- Single execution of the query
- "Compatible" with EMF

- Gremlin metamodel (around 100 classes)
- ATL Transformation
 - OCL-to-Gremlin mapping
 - Query composition
 - 70 rules and helpers
- Customized Gremlin engine
- Model element reification mechanism

Mogwai: load a Mogwai resource

1. Register the BlueprintsPersistenceBackendFactory.
2. Create a ResourceSet and register the PersistentResourceFactory
3. Create a new MogwaiURI to locate a file-based resource.
4. Create and cast the resource.
5. Use the Mogwai API

```
PersistenceBackendFactoryRegistry.register(  
    MogwaiURI.MOGWAI__SCHEME, BlueprintsPersistenceBackendFactory.getInstance());
```

```
ResourceSet resourceSet = new ResourceSetImpl();  
resourceSet.getResourceFactoryRegistry().getProtocolToFactoryMap()  
    .put(MogwaiURI.MOGWAI__SCHEME, MogwaiResourceFactory.getInstance());
```

```
URI uri = MogwaiURI.createMogwaiURI(new File("<db_path>"));
```

```
MogwaiResource resource = (MogwaiResource) resourceSet.createResource(uri);  
resource.load(Collections.emptyMap());
```

```
// Use EMF Resource with enhanced querying API  
resource.query([...]);
```

Mogwai: load and execute an OCL query

1. Create a MogwaiQuery
2. Query the resource and get a QueryResult
3. Retrieve the database results, execution time, executed query ...

```
MogwaiQuery query = OCLQueryBuilder.newBuilder()  
    .fromURI(URI.createURI("ocl/singletonMethods.ocl"))  
    .build();
```

```
QueryResult result = resource.query(MogwaiQuery);  
result.isSingleResult(); // returns only one element?  
result.resultSize(); // number of results  
result.getExecutedQuery(); // get the executed database query  
result.getComputationTime(); // time to compute the query  
result.getResults(); // Collection<Object> of database results
```

Mogwai: manipulate query results

1. Get a NeoEMFQueryResult
2. Reify the results (if possible¹)
3. Navigate your model elements using the standard EMF API

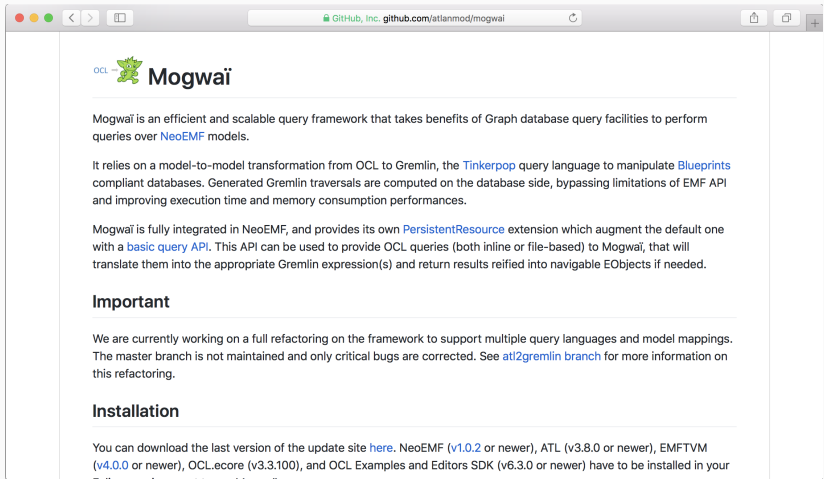
```
MogwaiQuery query = OCLQueryBuilder.newBuilder()
    .fromURI(URI.createURI("ocl/singletonMethods.ocl"))
    .build();
```

```
NeoEMFQueryResult result = resource.query(MogwaiQuery);
if(result.isReifiable()) {
    List<EObject> eObjects = result.reifyResults();
    for(EObject e : eObjects) {
        System.out.println(((MethodDeclaration)e).getName());
    }
}
```

¹Primitive types cannot be reified

- ModelDatastore abstraction
 - Support for different data stores
 - Easily extensible
 - Generic queries
- Prototype support for model transformations (Gremlin-ATL [DJSC17])
- Data migration operations
- Large model validation (presenting this work at 15:00 at OCL'18)

Mogwai: project website



- `https://github.com/atlanmod/mogwai`
- Open source project under the Eclipse Public License 2.0

NeoEMF

- Select the NoSQL database adapted to a modeling scenario
- Transparent EMF integration
- On-demand loading

Mogwai

- Benefit from the capabilities of NeoEMF/Graph backend
- Translates OCL queries into Gremlin traversals
- Bypasses low-level modeling APIs



Gwendal Daniel, Frédéric Jouault, Gerson Sunyé, and Jordi Cabot.

Gremlin-ATL: a scalable model transformation framework.

In *Automated Software Engineering (ASE), 2017 32nd IEEE/ACM International Conference on*, pages 462–472. IEEE, 2017.



Gwendal Daniel, Gerson Sunyé, Amine Benelallam, Massimo Tisi, Yoann Vernageau, Abel Gómez, and Jordi Cabot.

NeoEMF: a multi-database model persistence framework for very large models.

In *Science of Computer Programming*, pages 1–7. Elsevier, 2017.



Gwendal Daniel, Gerson Sunyé, and Jordi Cabot.

Mogwaï: A framework to handle complex queries on large models.

In *Tenth IEEE International Conference on Research Challenges in Information Science, RCIS 2016, Grenoble, France, June 1-3, 2016*, pages 1–12. IEEE, 2016.