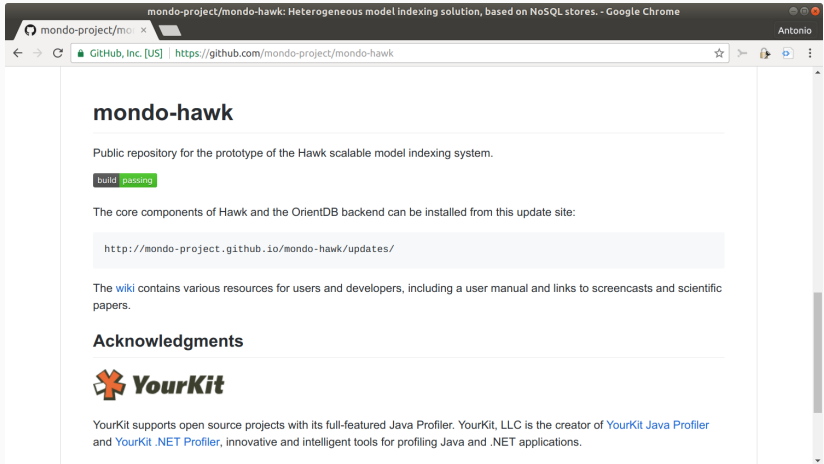


Taming Large Models with Hawk and NeoEMF

A. García-Domínguez, D. S. Kolovos, K. Barmpis, G. Daniel, G. Sunyé
MoDELS'2018, 14–19 October 2018


Hawk

Hawk: project website




- <https://github.com/mondo-project/mondo-hawk>
- Recently accepted as Eclipse Incubator project (moving soon!)

Hawk: deployment



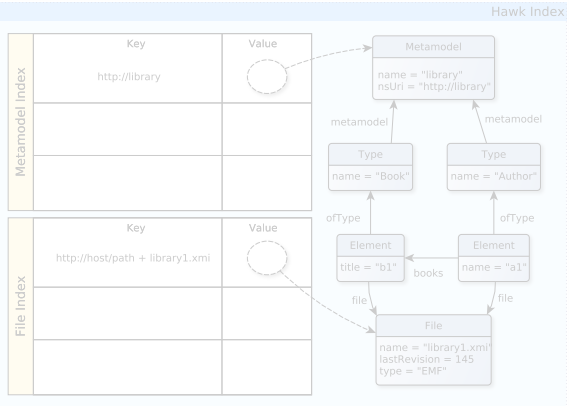

- Hawk can run as Eclipse plug-in, Java library, or network service
- We can have it watch over various types of locations:
 - Version control systems (SVN/Git repositories)
 - File stores (local folders, Eclipse workspaces, HTTP locations)

Hawk: component-based architecture



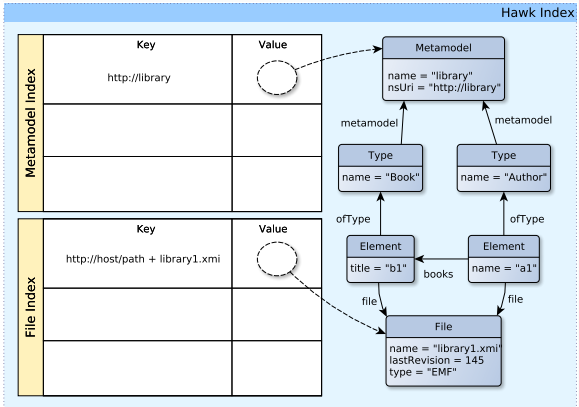

- Core: incremental graph updating + component interfaces
- Backends: Neo4j (fastest), OrientDB (multi-master), Greycat
- Clients: Eclipse GUI, cross-language Apache Thrift web services
- Query engines: Epsilon Object/Pattern Languages, OrientDB SQL
- Model parsers: EMF/Modelio models, Eclipse plug-in manifests...

Hawk: example for a library model




We go from these model files...

Hawk: example for a library model



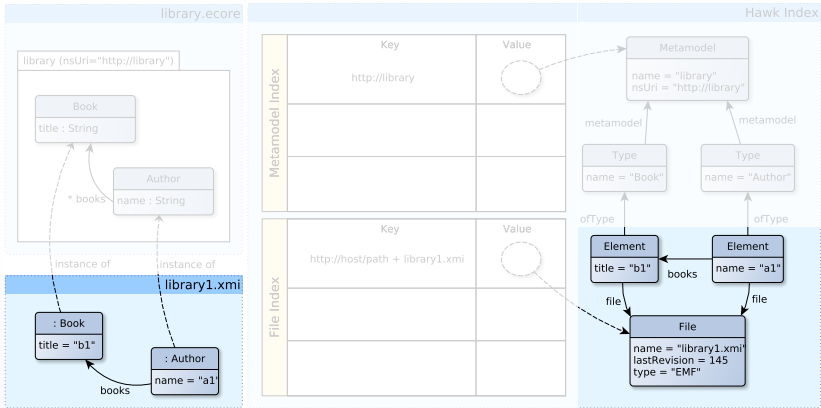
... to these NoSQL graphs.

Hawk: example for a library model




- Ecore packages → metamodel nodes
- Ecore classes → type nodes

Hawk: example for a library model



- Physical files → file nodes
- Model elements → element nodes

Hawk: example for a library model



- MM index: package URI → metamodel node
- File index: file path → file node
- Users can define custom indices by attribute/expression

Demo time!

Let's index a Java model and find
singletons.

Hawk: indexed attributes

Finding a type by its name

```
return TypeDeclaration.all.selectOne(td |  
    td.name.identifier = 'IdInputFileObject');
```

This normally involves...

1. Iterating over all types
2. Following the “name” reference
3. Comparing the name

Hawk can replace this with a lookup

We only need to tell it to index “SimpleName.identifier”:

```
return SimpleName.all  
    .select(sn | sn.identifier='IdInputFileObject')  
    .eContainer.select(c|c.isKindOf(TypeDeclaration));
```

Hawk: derived attributes

Original query for finding singletons

```
return TypeDeclaration.all.select(td | td.bodyDeclarations.exists(  
  md:MethodDeclaration | md.returnType.isTypeOf(SimpleType)  
  and md.returnType.name.fullyQualifiedName  
    = td.name.fullyQualifiedName  
  and md.modifiers.exists(mod:Modifier | mod.public = true)  
  and md.modifiers.exists(mod:Modifier | mod.static = true)  
));
```

Can we do it faster?

- Checking if a method is public or static requires traversing references
- Same goes for checking if it returns an instance of itself
- In Hawk, **we can precompute this**
- When files change, only the affected values are recomputed

Hawk: use of derived attributes as precomputed values

Original query

```
return TypeDeclaration.all.select(td | td.bodyDeclarations.exists(  
  md:MethodDeclaration | md.returnType.isTypeOf(SimpleType)  
  and md.returnType.name.fullyQualifiedName  
    = td.name.fullyQualifiedName  
  and md.modifiers.exists(mod:Modifier | mod.public = true)  
  and md.modifiers.exists(mod:Modifier | mod.static = true)  
));
```

Changed to use derived attributes on MethodDeclaration

```
return TypeDeclaration.all.select(td |  
  td.bodyDeclarations.exists(md:MethodDeclaration |  
    md.isPublic = true  
    and md.isStatic = true  
    and md.isSameReturnType = true));
```

Hawk: derived attributes are also indexed

Revised query

```
return TypeDeclaration.all.select(td|
  td.bodyDeclarations.exists(md:MethodDeclaration |
    md.isPublic = true
    and md.isStatic = true
    and md.isSameReturnType = true));
```

Can we do it faster?

- Right now, we need to go through *all* type declarations and then filter by methods
- What if we go from the methods to the types instead?
- In Hawk, **top-level selects can replace iteration with lookups when using derived attributes**

Hawk: use of derived attributes as index keys

Previous query

```
return TypeDeclaration.all.select(td |  
  td.bodyDeclarations.exists(md:MethodDeclaration |  
    md.isPublic = true  
    and md.isStatic = true  
    and md.isSameReturnType = true));
```

Revised to use index, by using derived attributes at the top level

```
return MethodDeclaration.all.select(md |  
  md.isPublic = true and md.isStatic = true  
  and md.isSameReturnType = true  
).collect( td | td.eContainer ).asSet;
```

Hawk: flagging singletons directly

Previous query

```
return MethodDeclaration.all.select(md |  
    md.isPublic = true and md.isStatic = true  
    and md.isSameReturnType = true  
)collect( td | td.eContainer ).asSet;
```

Can we do it faster?

- We could just flag types that are singletons
- This derived attribute might be less reusable, however

Hawk: final query for finding singletons

Previous query

```
return MethodDeclaration.all.select(md |  
    md.isPublic = true and md.isStatic = true  
    and md.isSameReturnType = true  
).collect( td | td.eContainer ).asSet;
```

Final query

```
return TypeDeclaration.all.select(td | td.isSingleton = true);
```

Demo time!

This time, we will show how to use indexed
and derived attributes.

Derived edges

Toy example: Person metamodel

- Person metamodel, with “parents” references.
- We want to be able to quickly find siblings, grandparents, uncles/aunts, cousins, second-cousins, ancestors...
- We can precompute this in Hawk with **derived edges**

Derivation logic for “grandparents”

We need a flat list and not a list of lists, so we use “flatten”:

```
return self.parents.parents.flatten;
```

Derivation logic for “siblings”


We can travel references in reverse with “revRefNav_name”:

```
return self.parents.revRefNav_parents.flatten.excluding(self);
```


Last demo for Hawk.

We will show derived edges this time.

Hawk: integration into SOFTEAM Constellation [GDBK⁺16]



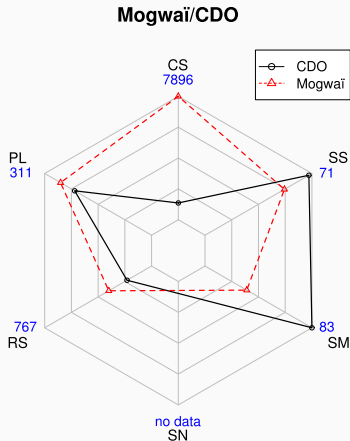

**Indexing times and index sizes
(OrientDB backend)**



**Code generation times:
Modelio (MT), Hawk
(HT)**

- Constellation: collaboration platform over Modelio models
- SOFTEAM needed search, couldn't change persistence
- Integrated Hawk as a library: initial indexing cost quickly paid off

Hawk: stress-testing remote query APIs [GDBK⁺17]



- Included CDO, Hawk, Mogwai and ranged 1–64 clients
- Reverse reference navigation was crucial in having the SN Train Benchmark [SIRV17] query run quickly

Hawk: time-aware queries over versioned models [GDBP18]

```
var rs = RewardTableRow.latest.all.collect(row | row.getRewardShifts()).flatten();  
return Sequence { rs.min(), rs.max(), rs.average() };
```

```
operation RewardTableRow getRewardShifts(): Sequence {  
  var v = self.versions;  
  if (v.size <= 1) { return Sequence {}; }  
  else { return v.subList(0, v.size - 1).collect(v | v.value - v.prev.value); }  
}  
operation Sequence average() { return self.sum() / self.size(); }
```

- Extended Hawk with Greycat temporal graph support and time-aware query engine / updater components
- Can index entire Subversion-based history of a model, and ask things about its history through a new set of time-aware primitives
- Above query is finding descriptive stats for reward table shifts in a models@run.time system
- Presenting this work at 14:00 (MRT'18) — hope to see you there!

Hawk: summing up

So far...

- Hawk is good for indexing an existing collection of model files
- You can efficiently answer queries from the index
- Indexed/derived features can be used to speed up queries

Ideas in the roadmap

- Extensible UI, covering differences in options for components
- Horizontal scaling (a flock of Hawks?)
- Web UI based on Thrift API
- More backends! (Triple stores? Neo4j 3.x? MapDB?)
- Better Git connector (JGit-based)
- Extending EOL with Linear Temporal Logic
- Visualizations based on time-aware queries



Antonio Garcia-Dominguez, Konstantinos Barmpis, Dimitrios S Kolovos, Marcos Aurelio Almeida da Silva, Antonin Abherve, and Alessandra Bagnato.

Integration of a graph-based model indexer in commercial modelling tools.

In *Proceedings of the ACM/IEEE 19th International Conference on Model Driven Engineering Languages and Systems*, pages 340–350, Saint Malo, France, 2016. ACM Press.



Antonio Garcia-Dominguez, Konstantinos Barmpis, Dimitrios S. Kolovos, Ran Wei, and Richard F. Paige.

Stress-testing remote model querying APIs for relational and graph-based stores.

Software & Systems Modeling, pages 1–29, June 2017.



A. García-Domínguez, N. Bencomo, and Luis H. García Paucar.
Reflecting on the past and the present with temporal graph-based models.

In Proceedings of the 13th International Workshop on Models@run.time, 2018.

To be published.



Gábor Szárnyas, Benedek Izsó, István Ráth, and Dániel Varró.
The Train Benchmark: cross-technology performance evaluation of continuous model queries.

Software & Systems Modeling, January 2017.